

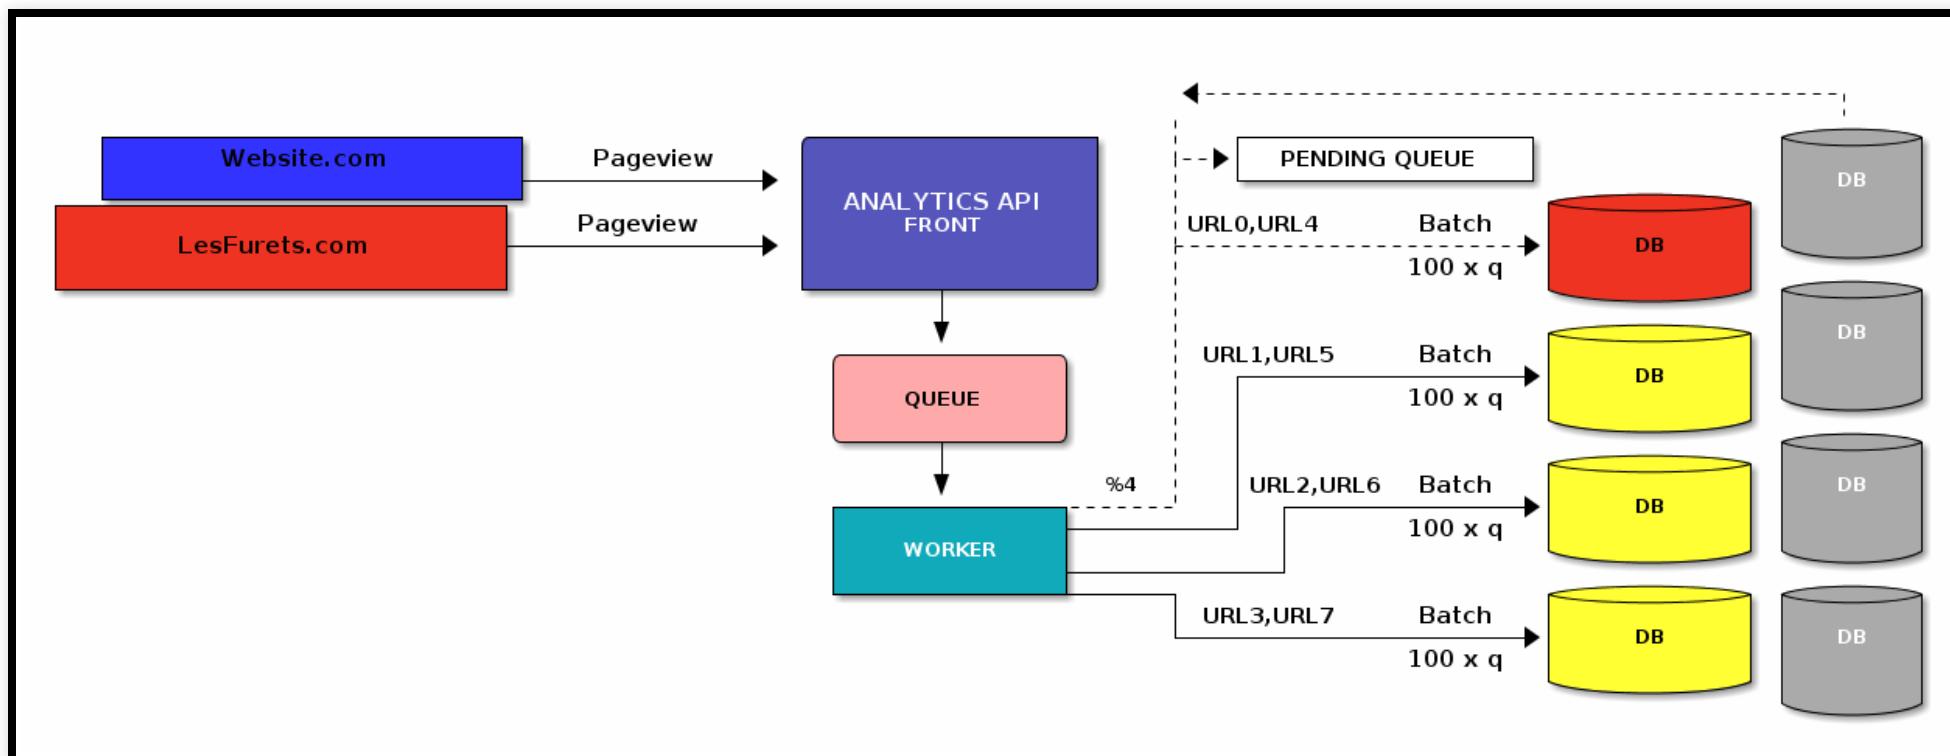
# Introduction to Apache Cassandra

Andrei Arion, LesFurets.com, tp-bigdata@lesfurets.com

# Plan

- Motivation
- Apache Cassandra
  - Partitioning and replication
  - Consistency
- Practice: Tune consistency in Apache Cassandra

# Motivation



- do I build new features for customers?
- or just dealing with reading/writing the data?

# What went wrong?

- A single server cannot take the load  $\Rightarrow$  solution / complexity
  - *Better database*
    - easy to add/remove nodes (**scalling**)
    - transparent data distribution (**auto-sharding**)
    - handle failures (**auto-replication**)

# Plan

- Motivation
- Apache Cassandra
  - Partitioning and replication
  - Consistency
- Practice: Tune consistency in Apache Cassandra

# Apache Cassandra

- started @Facebook inspired by *BigTable* model and *Amazon DynamoDB*
- 2008 Open Source Project
- **Datastax**: commercial offering Datastax Enterprise
  - monitoring(OpsCenter) automating repairs backup...
  - other features: search, analytics, graph, encryption
- 2010 Top Level Apache Project
  - Datastax biggest committer
- 2016 Datastax/ASF disagreements([#Staxit](#))

# Apache Cassandra

***Open source, distributed database designed to handle large amounts of data across many commodity servers, providing high availability with no single point of failure. It offers robust support for clusters spanning multiple datacenters, with asynchronous masterless replication allowing low latency operations for all clients.***

# Cassandra

- column oriented NoSQL database
- distributed (data, query)
- resilient (no SPOF)
  - we can query any node ⇒ coordinator to dispatch and gather the results

# Cassandra terminology

RDBMS

Cassandra

---

Schema (set of tables)

***Keyspace***

---

Table

***Table/column family***

---

Row

Row

---

Database server

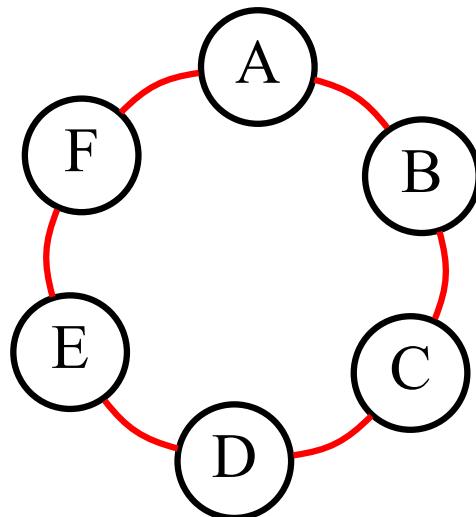
***Node***

---

Master/Slave

***Cluster: a set of nodes grouped in one or more datacenters (can span physical locations)***

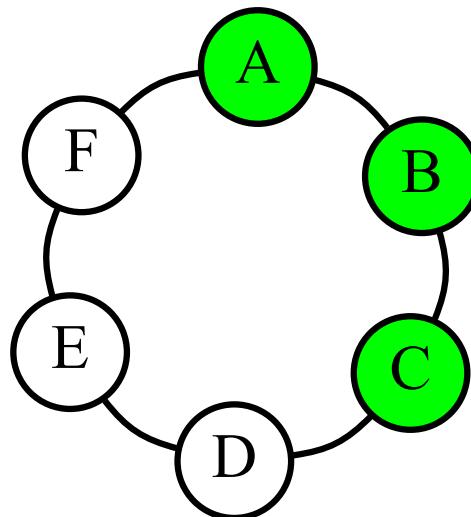
# Cassandra ring



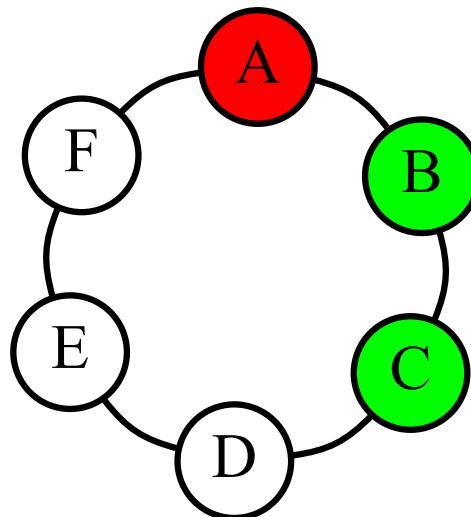
# Gossip

- peer-to-peer communication protocol
  - discover and share ***location*** and ***state information*** about nodes
  - persist gossip info locally to use when a node restarts
- **seed nodes** ⇒ bootstrapping the gossip process for new nodes joining the cluster

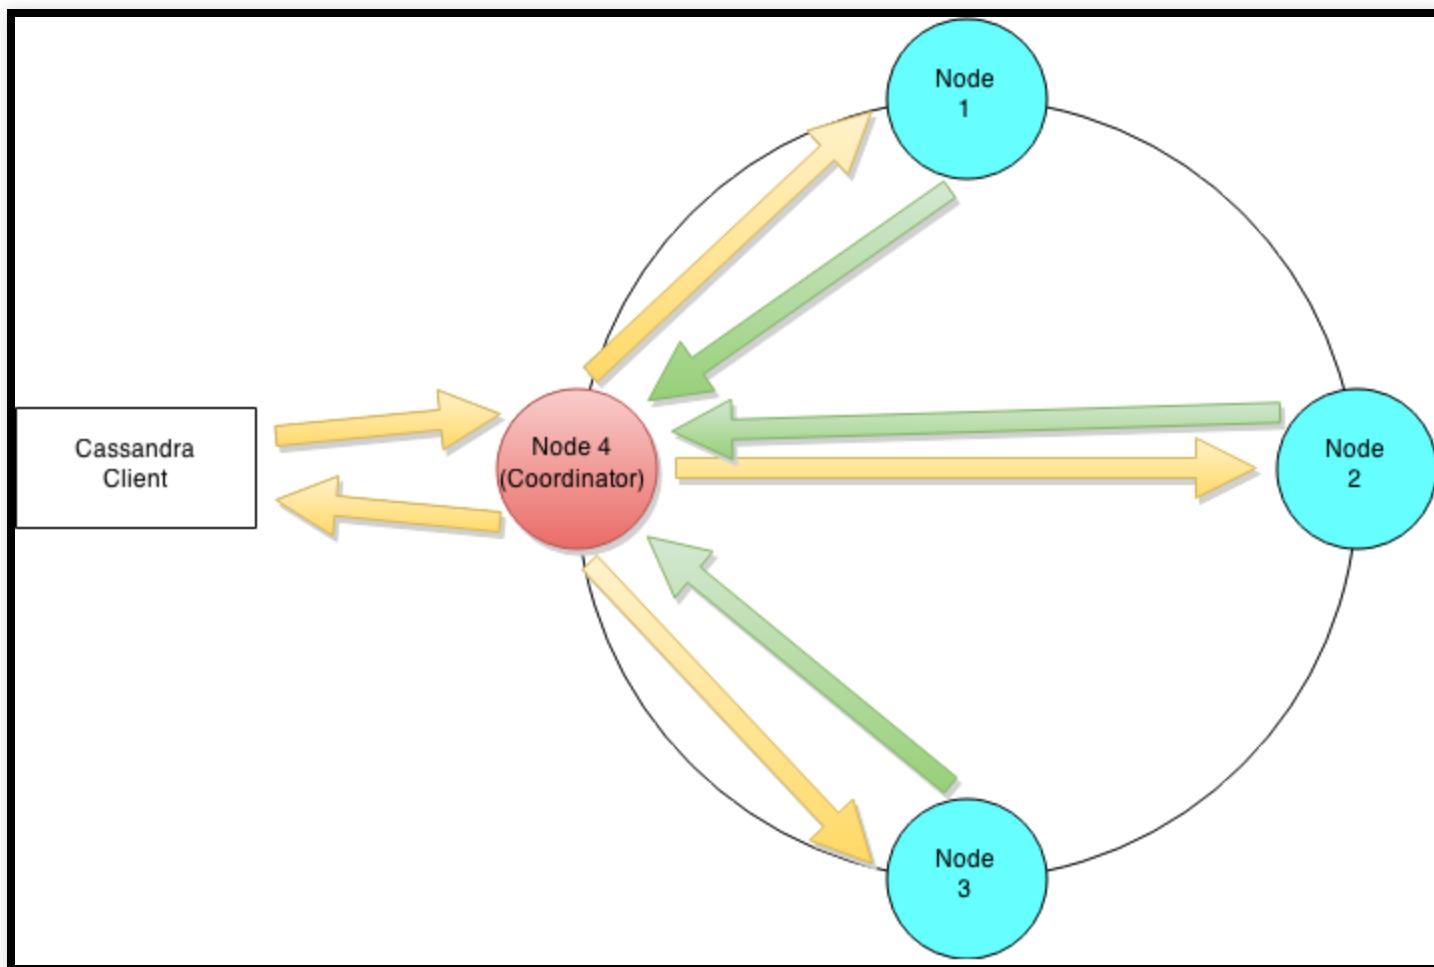
# Cassandra replicas



# Cassandra node failure



# Query



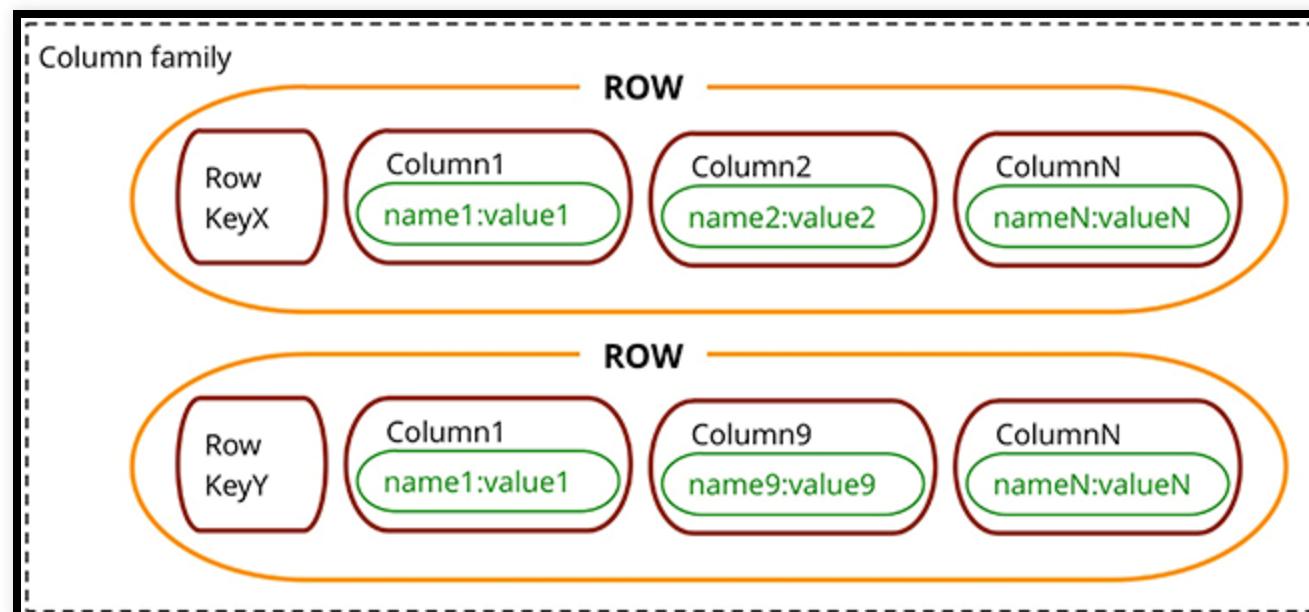
# Plan

- Motivation
- Apache Cassandra
  - Partitioning and replication
  - Consistency
- Practice: Tune consistency in Apache Cassandra

# Cassandra data model

- Stores data in **tables/column families** as rows that have many columns associated to a row key
- Free format: two rows can have different columns
- **Rows are the atomic aggregate** ⇒ distributed and replicated

**"Map<RowKey, SortedMap<ColumnKey, ColumnValue>>"**



# Data partitioning

- $C^*$  = single logical database spread across a cluster of nodes
- *How to divide data evenly around its cluster of nodes?*
- Problems:
  - distribute data **efficiently, evenly**
    - determining a node on which a specific piece of data should reside on
    - minimise the data movements when nodes join or leave the cluster
- Algorithm of **Consistent Hashing**

# Mapping data to nodes

- **Problem:** map  $k$  entries to  $n$  physical nodes
- *Naive hashing* ( $\text{NodeID} = \text{hash}(\text{key}) \% n$ )  $\Rightarrow$  remap a large number of keys when nodes join/leave the cluster
- *Consistent hashing*: only  $k/n$  keys need to be remapped on average

# Consistent Hashing

- Idea :
  - use a part of the data as a partition key
  - compute a hash value for each
  - The range of values from a consistent hashing algorithm is a fixed circular space which can be visualised as a ring.

# Partitioner

- **hash function** that derives a token from the primary key of a row
- determines which node will receive the *first replica*
- RandomPartitioner, Murmur3Partitioner, ByteOrdered

# Murmur3Partitionner

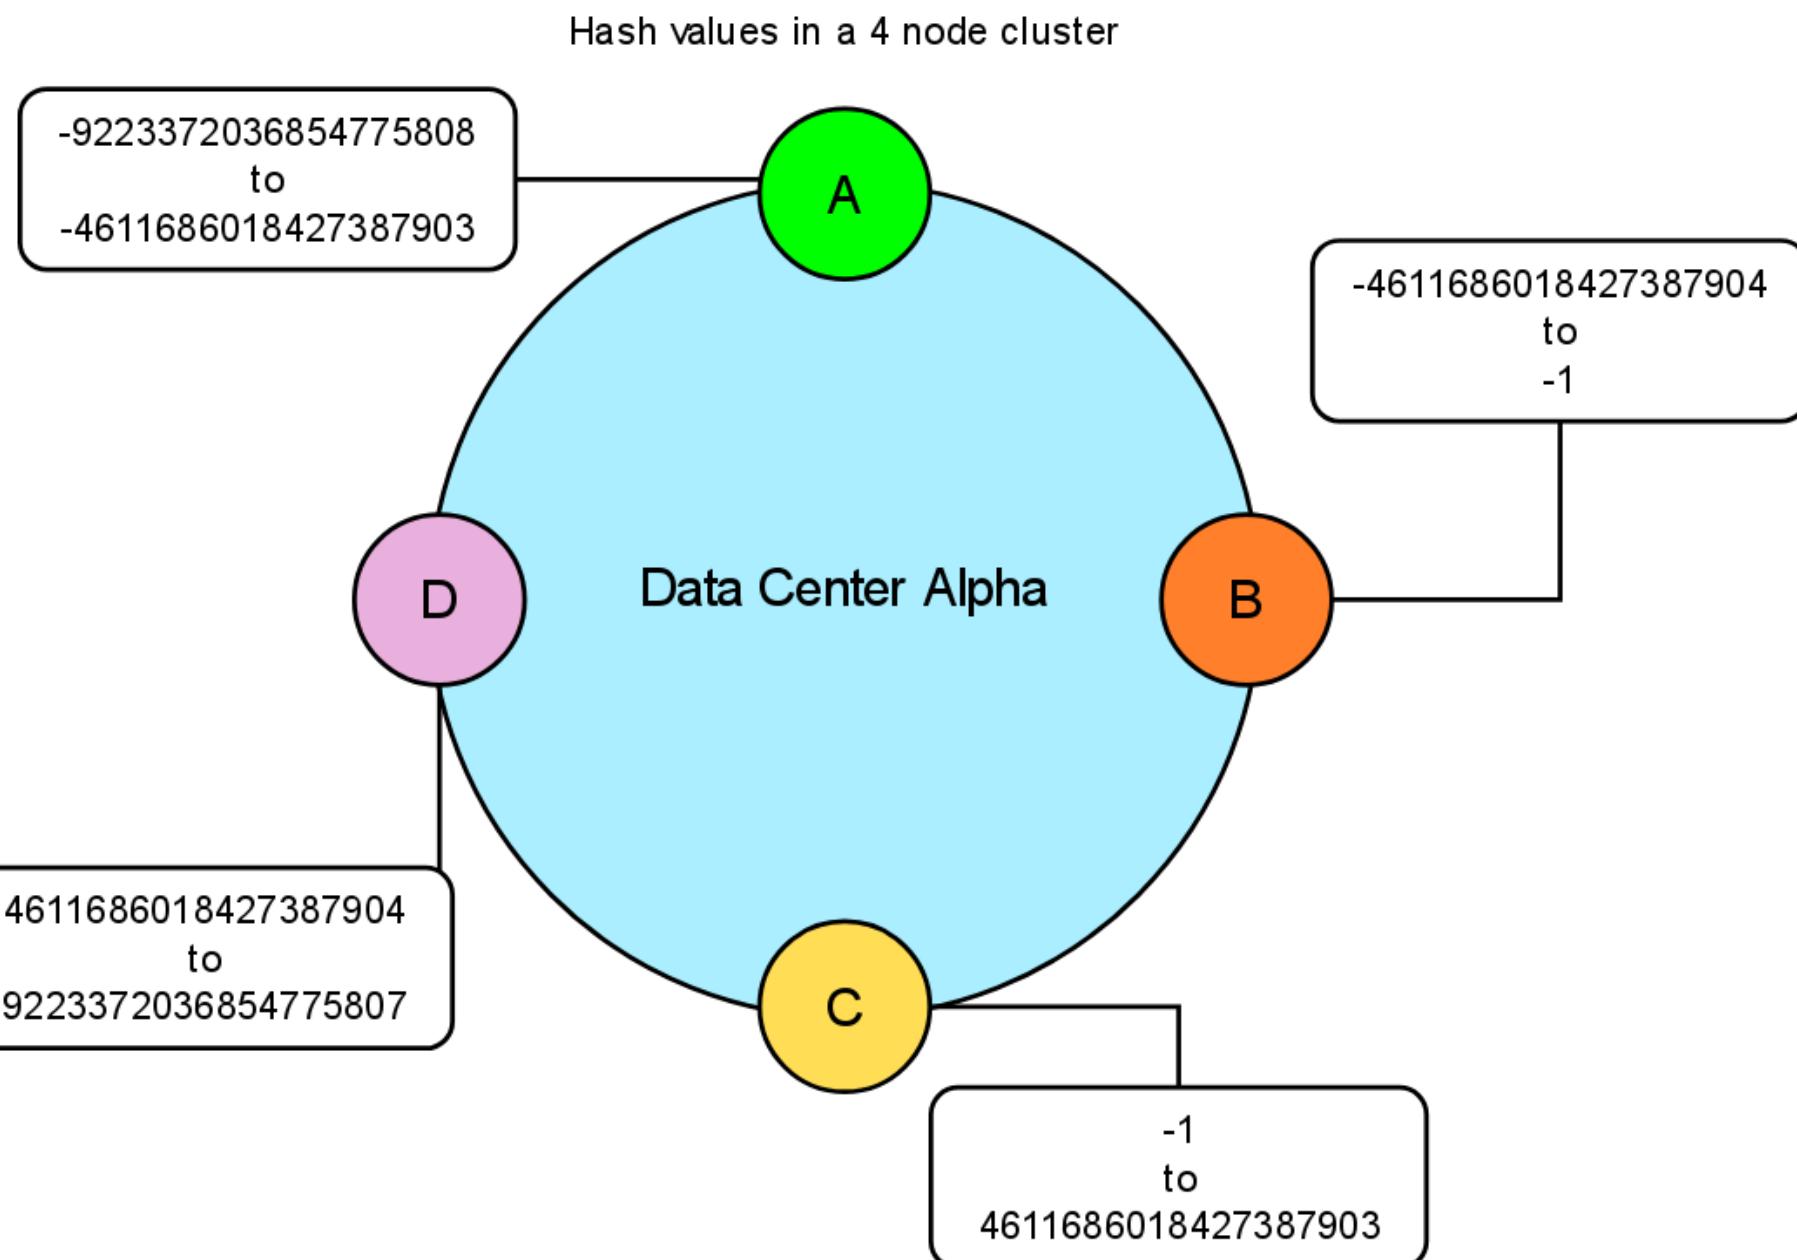
<b>name</b>	<b>age</b>	<b>car</b>	<b>gender</b>
jim	36	camaro	M
carol	37	bmw	F
johnny	12		M
suzy	10		F

Cassandra assigns a hash value to each partition key:

<b>Partition key</b>	<b>Murmur3 hash value</b>
jim	-2245462676723223822
carol	7723358927203680754
johnny	-6723372854036780875
suzy	1168604627387940318

# Consistent Hashing: mapping

Each node in the cluster is responsible for a range of data based on the hash value:



# Consistent Hashing: mapping

<b>No de</b>	<b>Start range</b>	<b>End range</b>	<b>Part itio n key</b>	<b>Hash value</b>
A	-92233720368547758 08	-461168601842738 7903	john ny	-67233728540367808 75
B	-46116860184273879 04	-1	jim	-22454626767232238 22
C	0	461168601842738 7903	suz y	116860462738794031 8
D	461168601842738790 4	922337203685477 5807	caro l	772335892720368075 4

# Data Replication

- create copies of the data, thus avoiding a single point of failure.
- **Replication Factor (RF)** = # of replica for each data
  - set at the *Keyspace* level

# How data is replicated: *Snitches*

- determine the *proximity of nodes within a ring*
  - *Dynamic snitch* ⇒ monitors the performance of reads from the replicas and chooses the best replica
  - *SimpleSnitch* ⇒ place the copy to the next available node (clockwise)
  - *RackInferringSnitch* ⇒ place copies of rows to different racks (same DC)

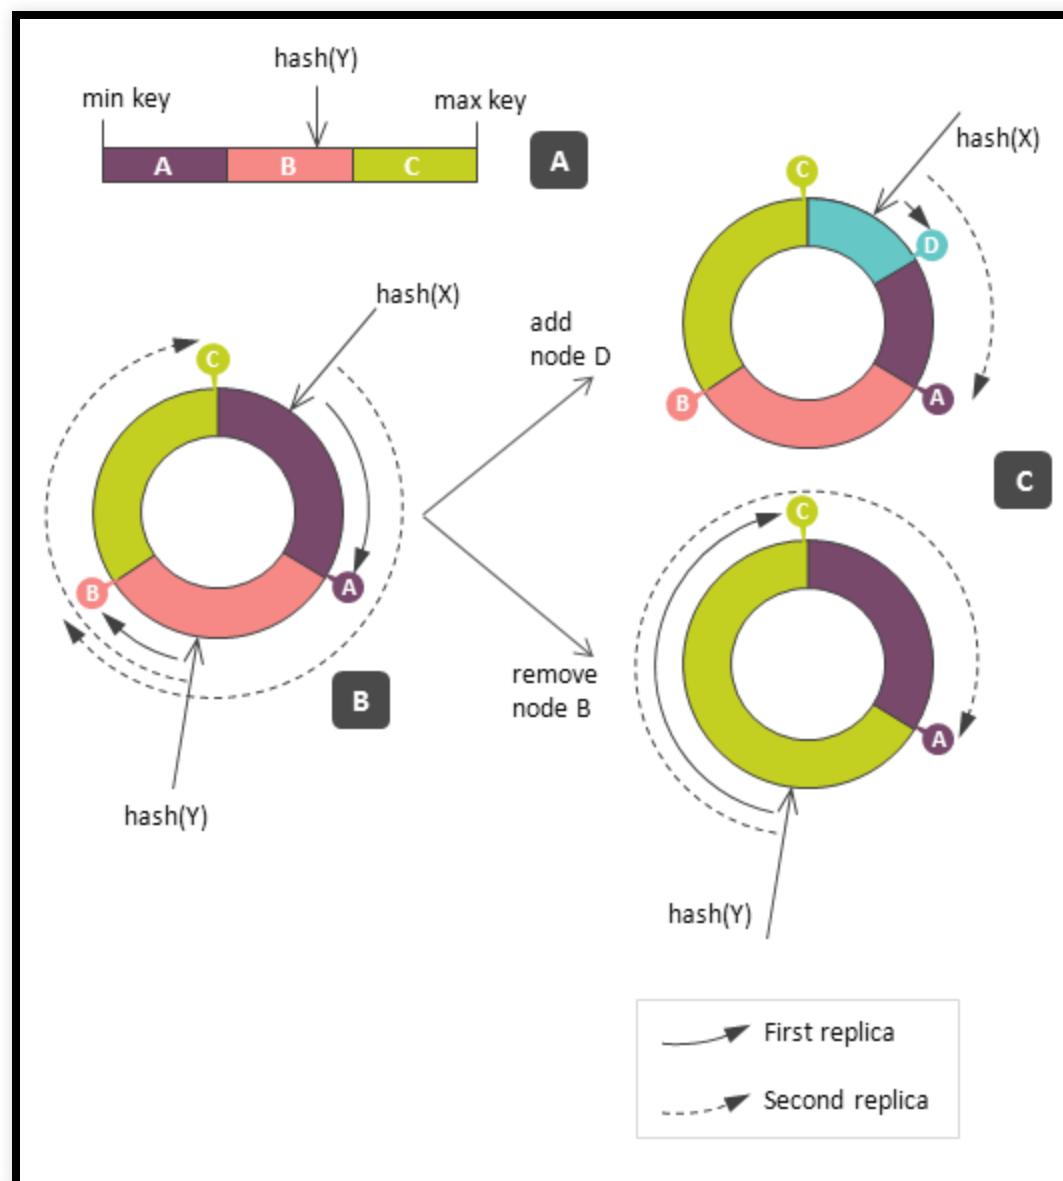
Configured in the `cassandra.yaml` file.

# How data is replicated: *Replication strategies*

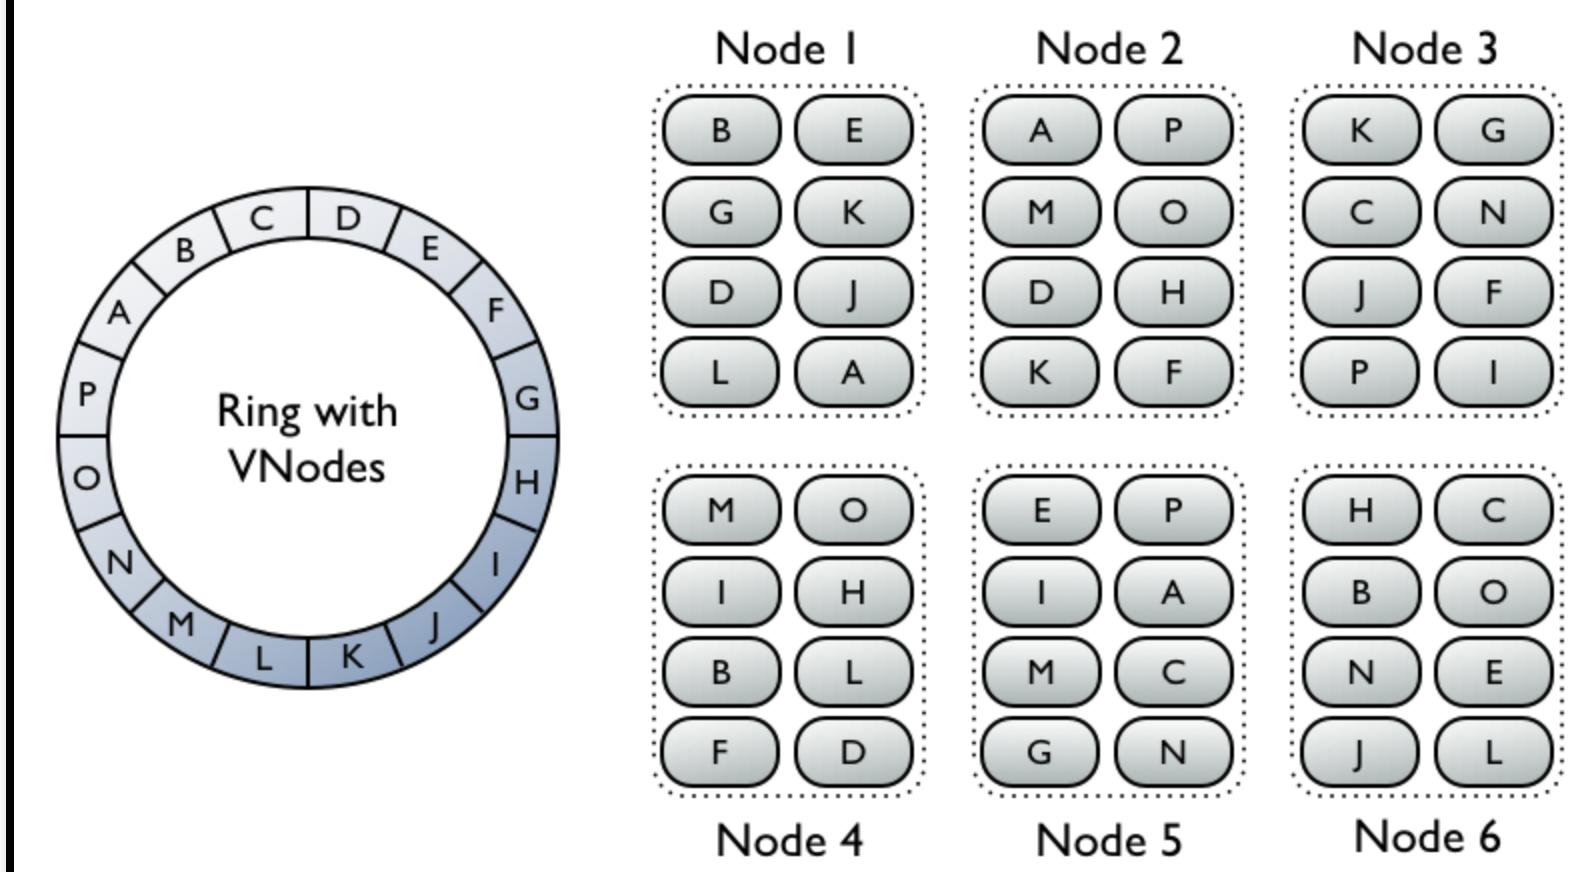
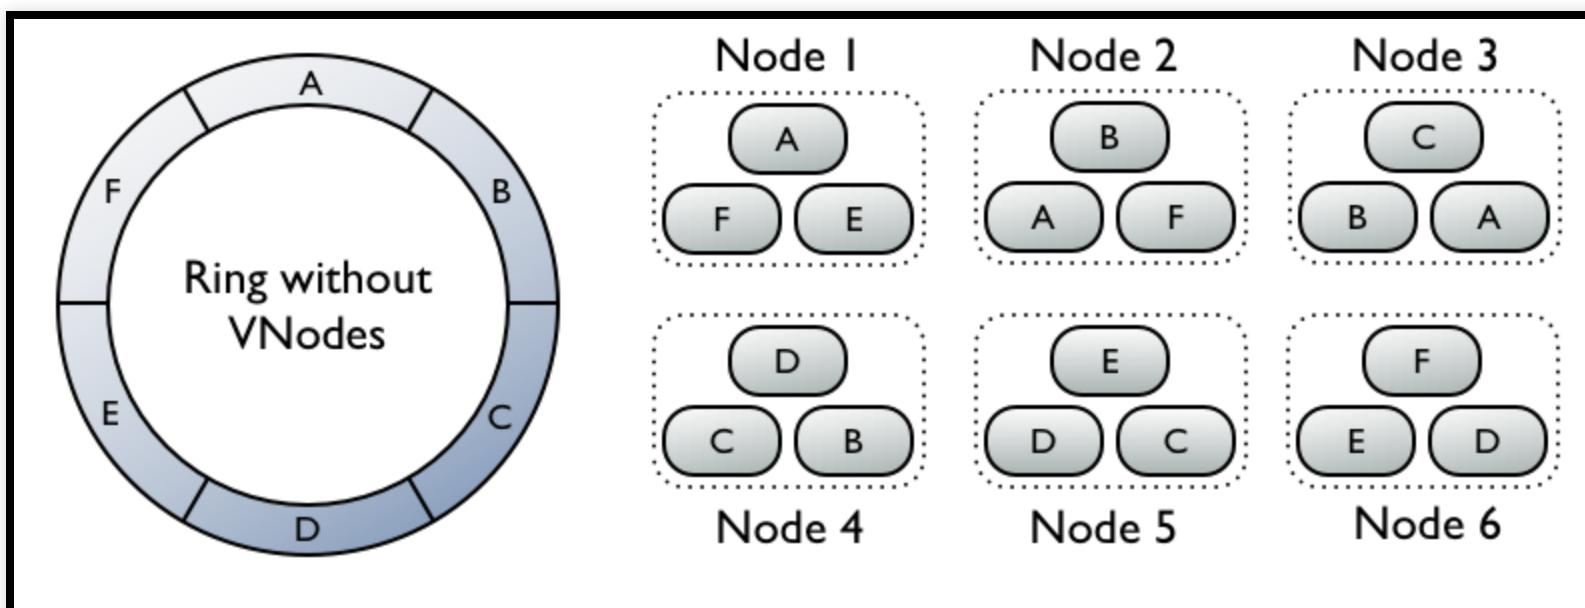
- use proximity information provided by snitches to determine locality of a copy
  - **SimpleStrategy**: use only for a single data center and one rack
  - **NetworkTopologyStrategy**: specifies how many replicas you want in each DC
- defined at **keyspace** level

```
CREATE KEYSPACE temperature
  WITH replication =
    {'class': 'SimpleTopologyStrategy', 'replication_factor':'2'};
```

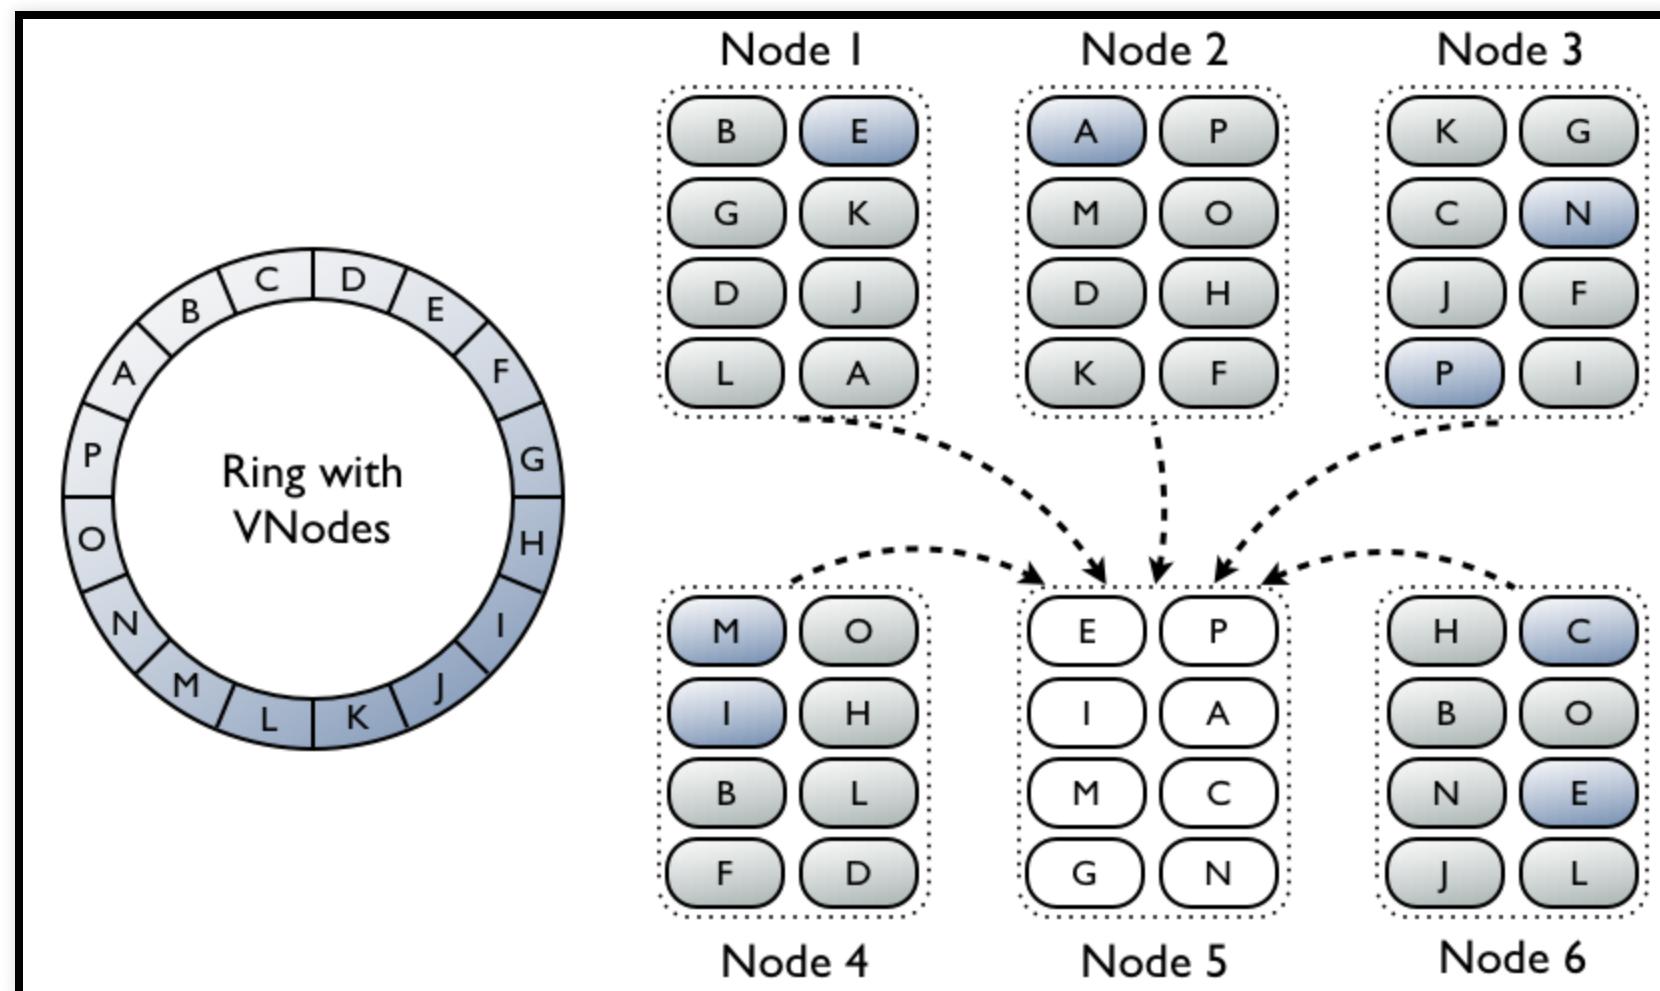
# Example: RF2 + SimpleSnitch + SimpleReplicationStrategy



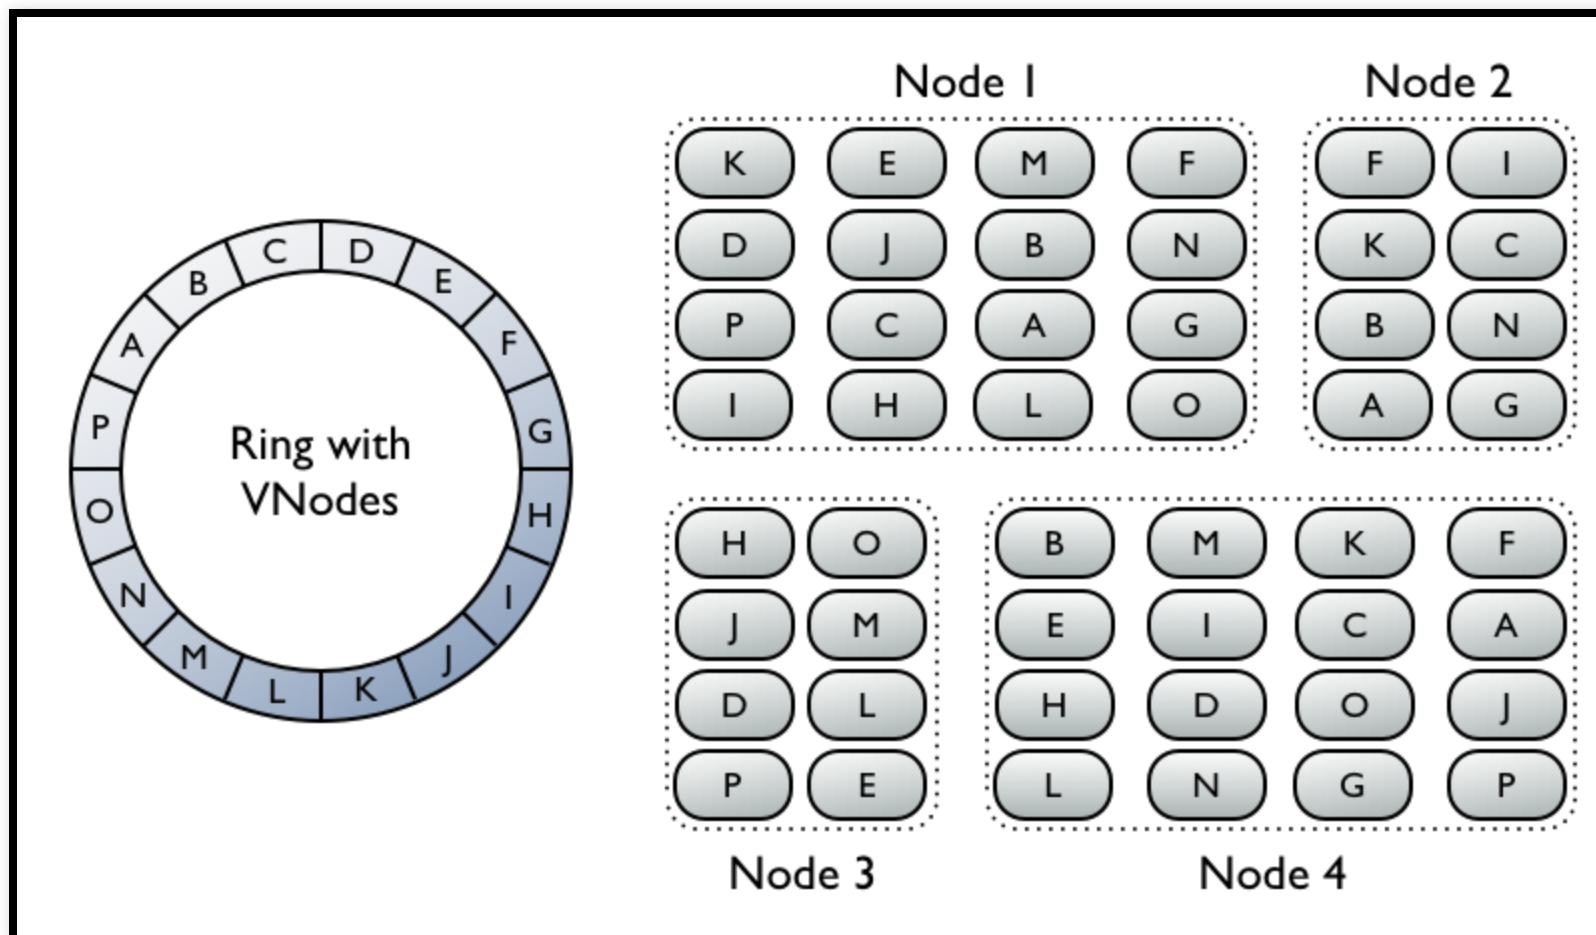
# Virtual nodes (VNodes)



# VNodes: remapping keys



# VNodes and heterogeneous clusters

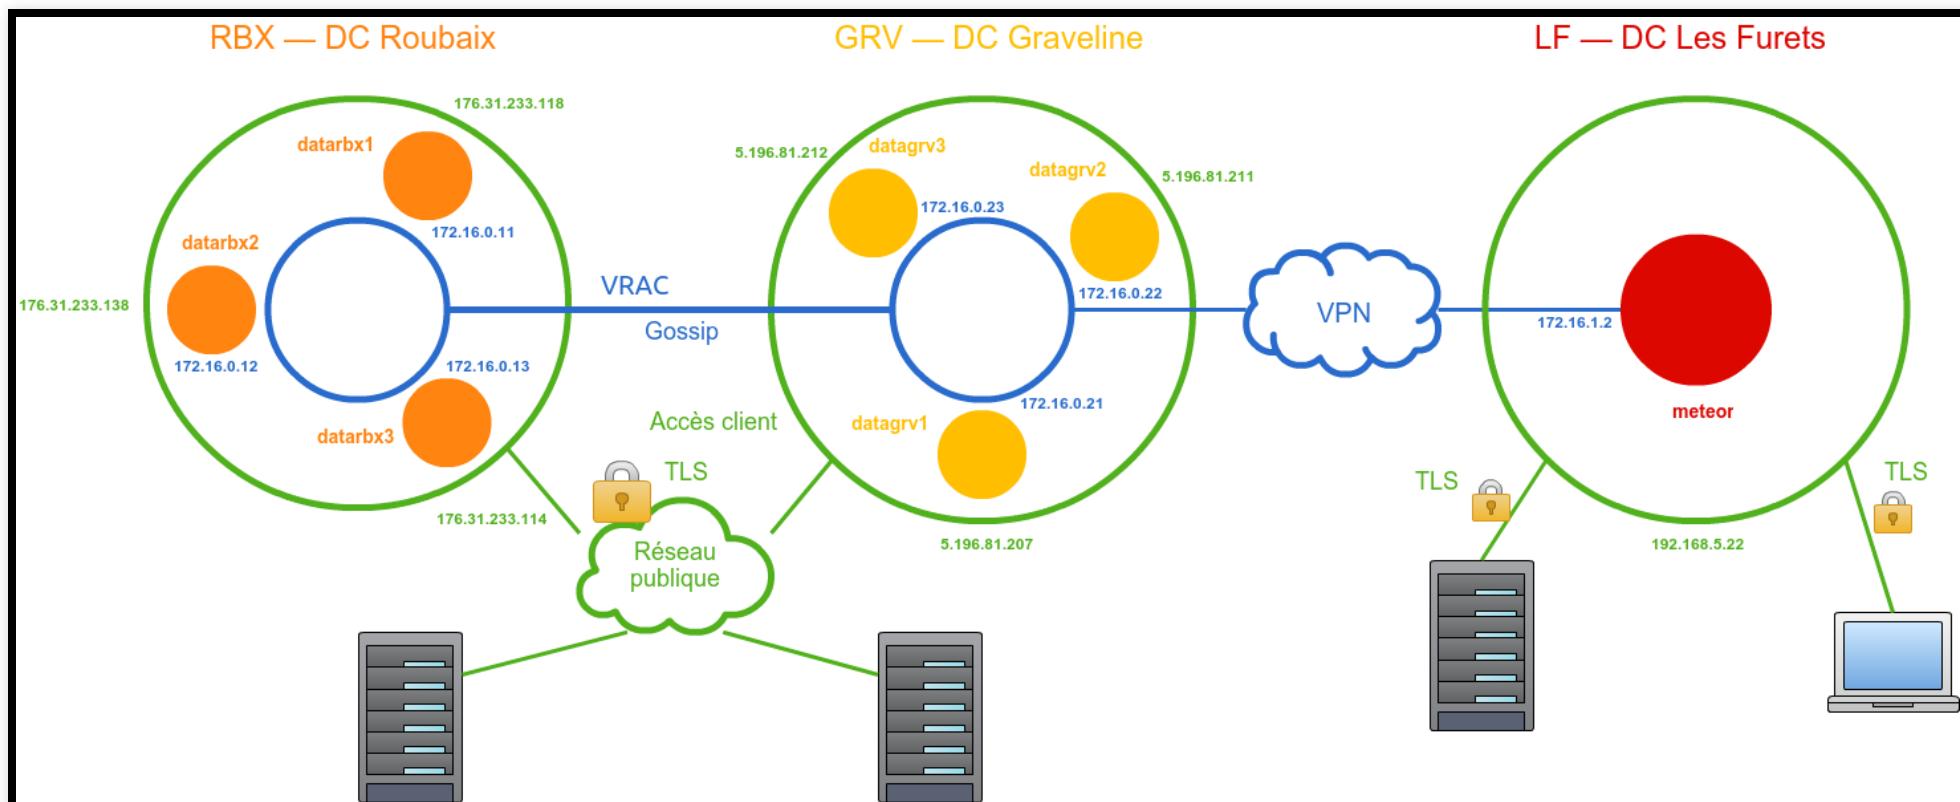


# Token allocation

- no VNODES  $\Rightarrow$  manual (initial-token="-29334..." dans `cassandra.yaml`)
  - need to be modified at each topology change
- VNODES ( `num_tokens` )
  - random slot allocation (< 3.0)
  - smart (3.0+)

# Example: NetworkTopologyStrategy

```
CREATE KEYSPACE lesfurets  
WITH replication =  
{'class': 'NetworkTopologyStrategy', 'RBX': 2, 'GRV': 2, 'LF': 1};
```



# Plan

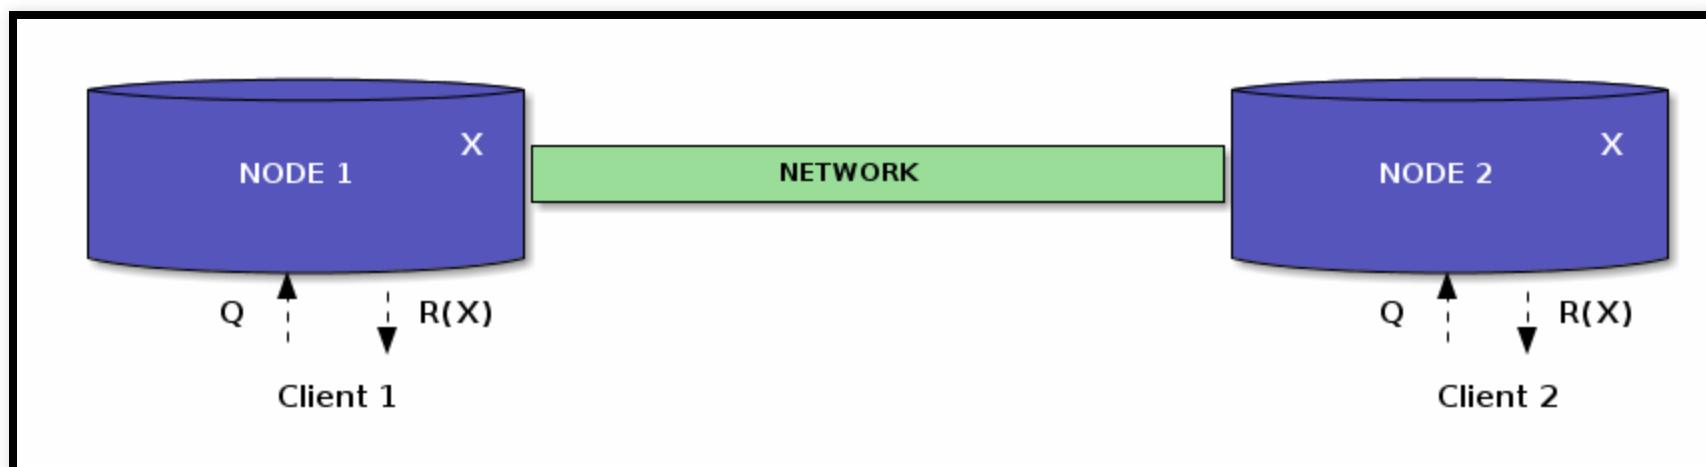
- Motivation
- Apache Cassandra
  - Partitioning and replication
  - **Consistency**
- Practice: Tune consistency in Apache Cassandra

# Properties of distributed systems

- **Consistency:** read is guaranteed to return the most recent write for a given client.
- **Availability:** non-failing node will return a reasonable response within a reasonable amount of time (no error or timeout)
- **Partition Tolerance:** the system will continue to function when network partitions occur.

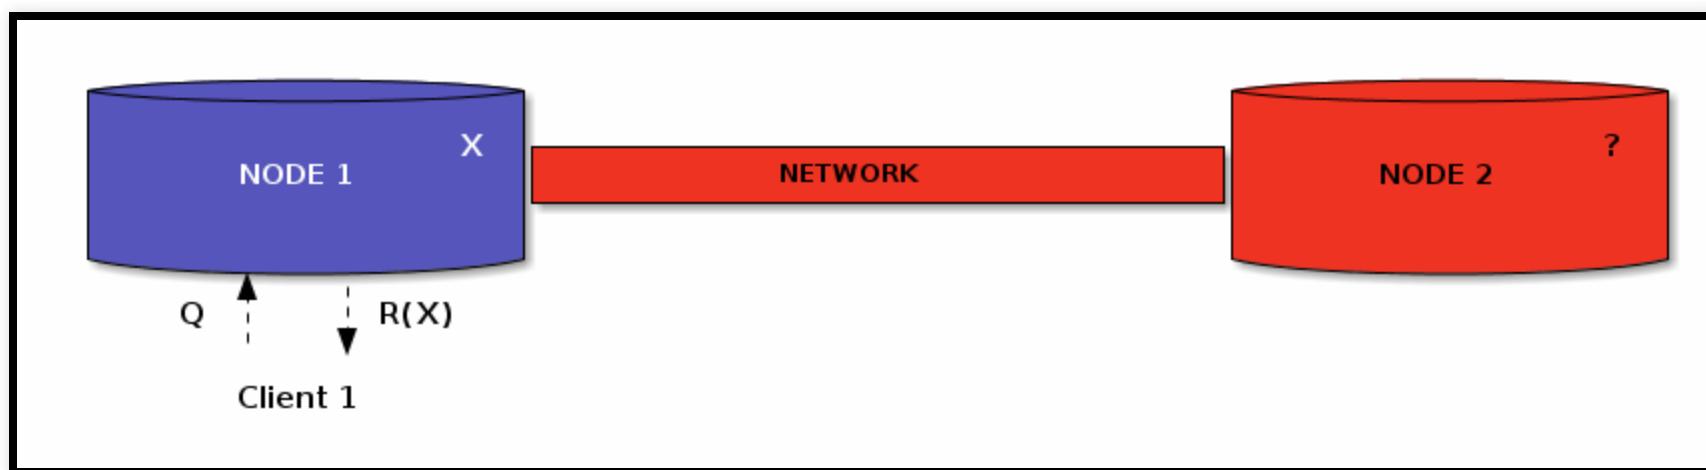
# Consistency

- *a read returns the most recent write*
- *eventually consistent* : guarantee that the system will evolve in a consistent state
  - provided there are no new updates, all nodes/replicas will eventually return the last updated value (~DNS)



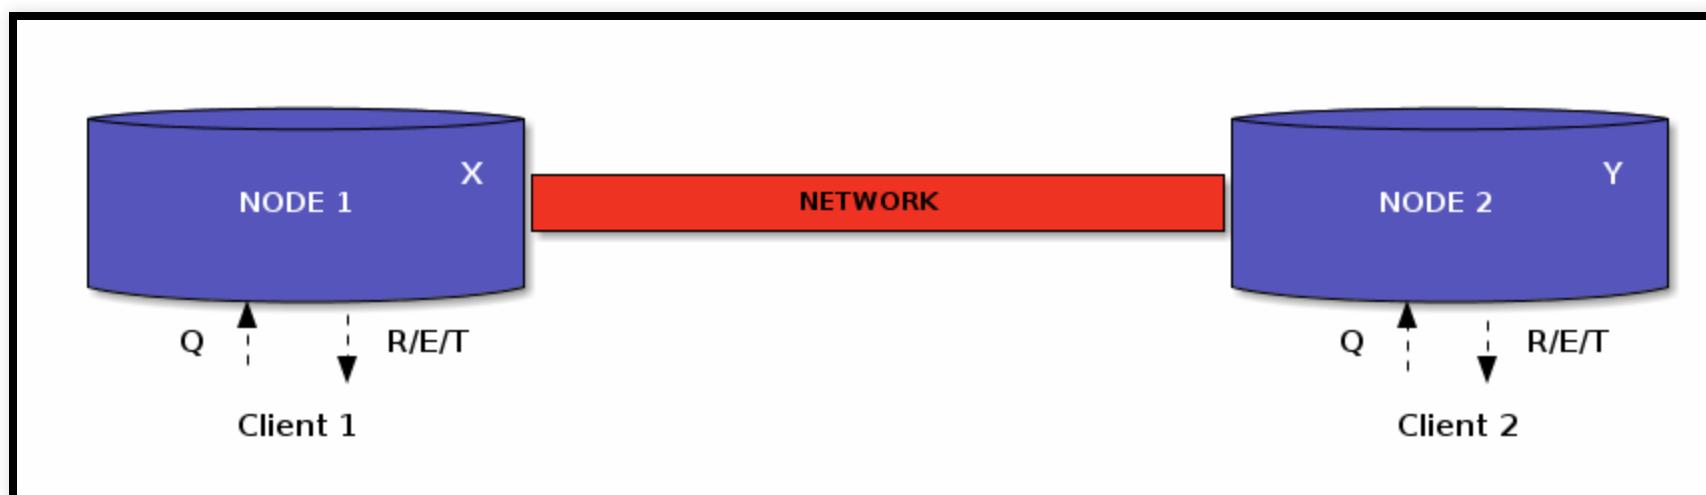
# Availability

- *a non-failing node will return a reasonable response (no error or timeout)*



# Partition tolerance

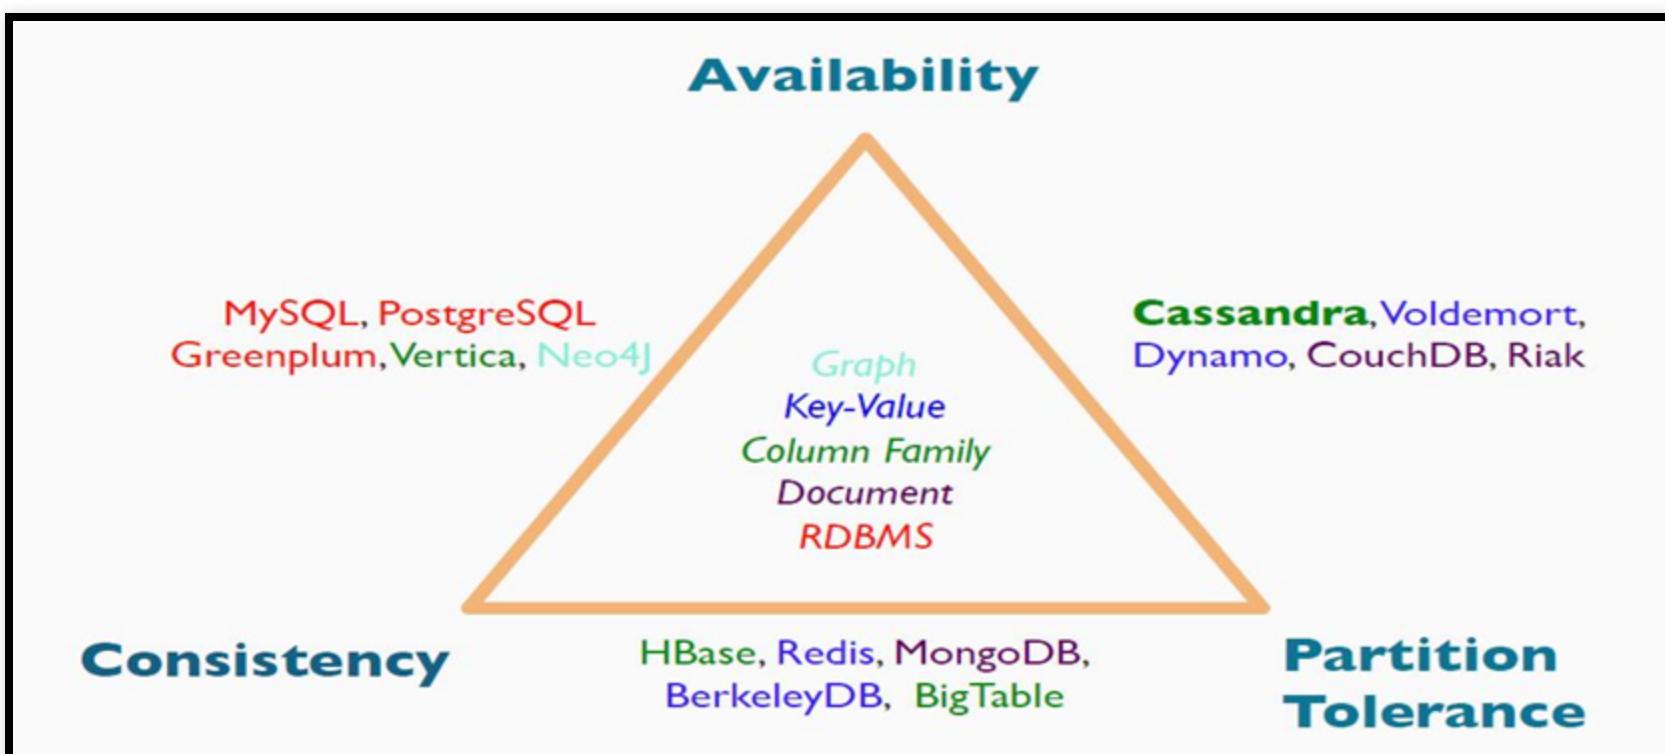
- *ability to function (return a response, error, timeout) when network partitions occur*



# Partition tolerance

- network **is** unreliable
- **you can choose** how to handle errors
  - return an old value
  - wait and eventually timeout, or return an error at once
- in practice: choice between AP and CP systems

# CAP theorem



# Cassandra consistency

- AP system
  - *eventually consistent*
    - without updates the system will converge to a consistent state due to *repairs*
  - *tunable consistency* :
    - Users can determine the consistency level by tuning it during **read** and **write** operations.

# Consistency Level (CL)

- mandatory **protocol-level** parameter for each query (read/write),
- #replicas in a cluster that must acknowledge the **read / write**
  - **write consistency - R:** #replicas on which the write *must succeed\** before returning an acknowledgment to the client application.
  - **read consistency - W:** #replicas that must *respond to a read* request before returning data to the client application
- default level: **ONE**
- most used: **ONE, QUORUM, ALL, ANY ... (LOCAL\_ONE, LOCAL\_QUORUM...)**

# Durability

- guarantees that writes, once completed, will survive permanently
  - appending writes to a commitlog first
    - default: flushed to disk every *commitlog\_sync\_period\_in\_ms*
    - batch mode ⇒ sync before ACK the write

# Which consistency level ?

W : Niveau de cohérence en écriture (nb de nœuds)

R : Niveau de cohérence en lecture (nb de nœuds)

RF : Facteur de réPLICATION

Cohérence immédiate si :

$$W + R > RF$$

W : ONE

R : ONE

Cohérence : à terme

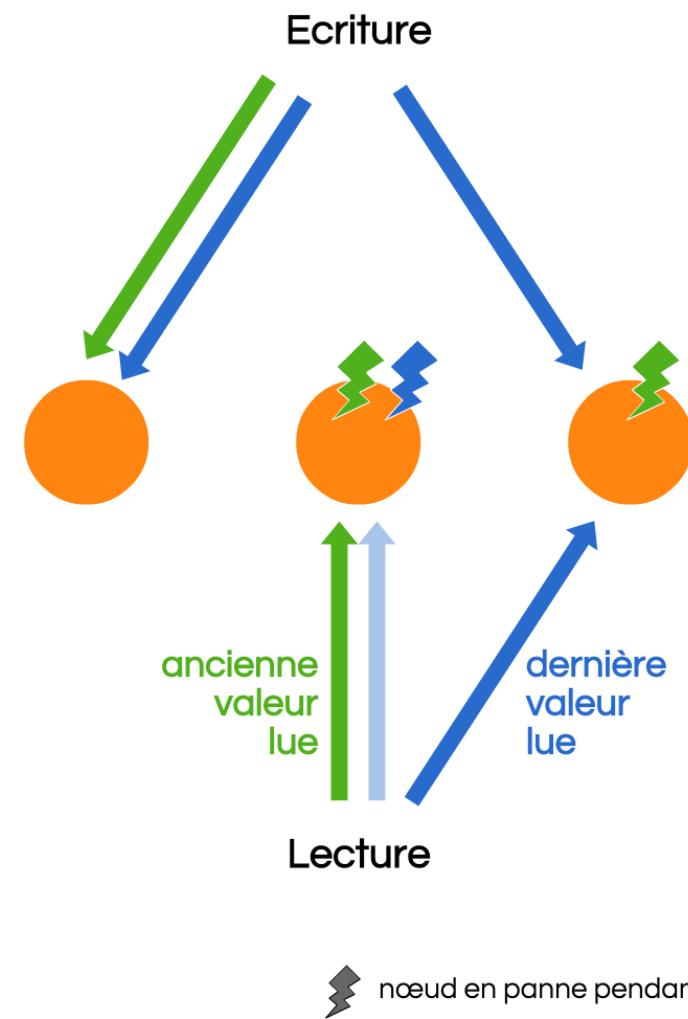
Disponibilité : 2 nœuds en panne

W : QUORUM

R : QUORUM

Cohérence : immédiate

Disponibilité : 1 nœud en panne



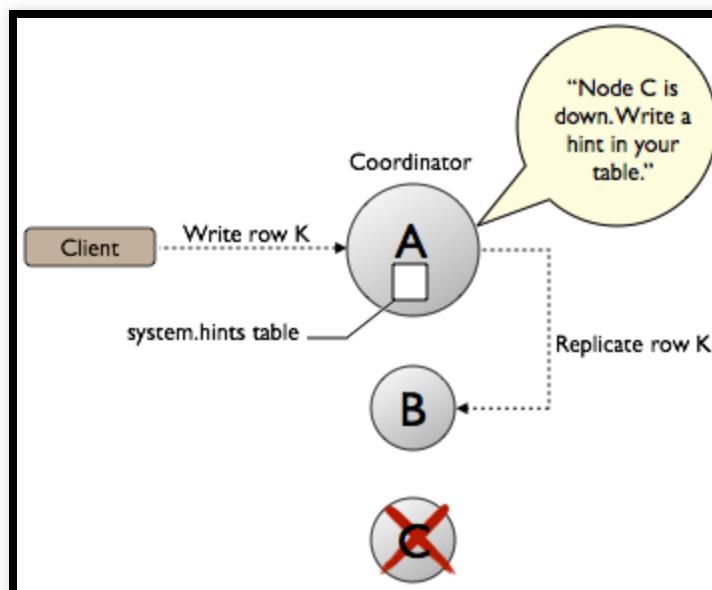
⚡ nœud en panne pendant l'écriture

# How to achieve consistency

1. hinted handoff
2. read repairs
3. anti-entropy repair (nodetool repair)

# Hinted handoff

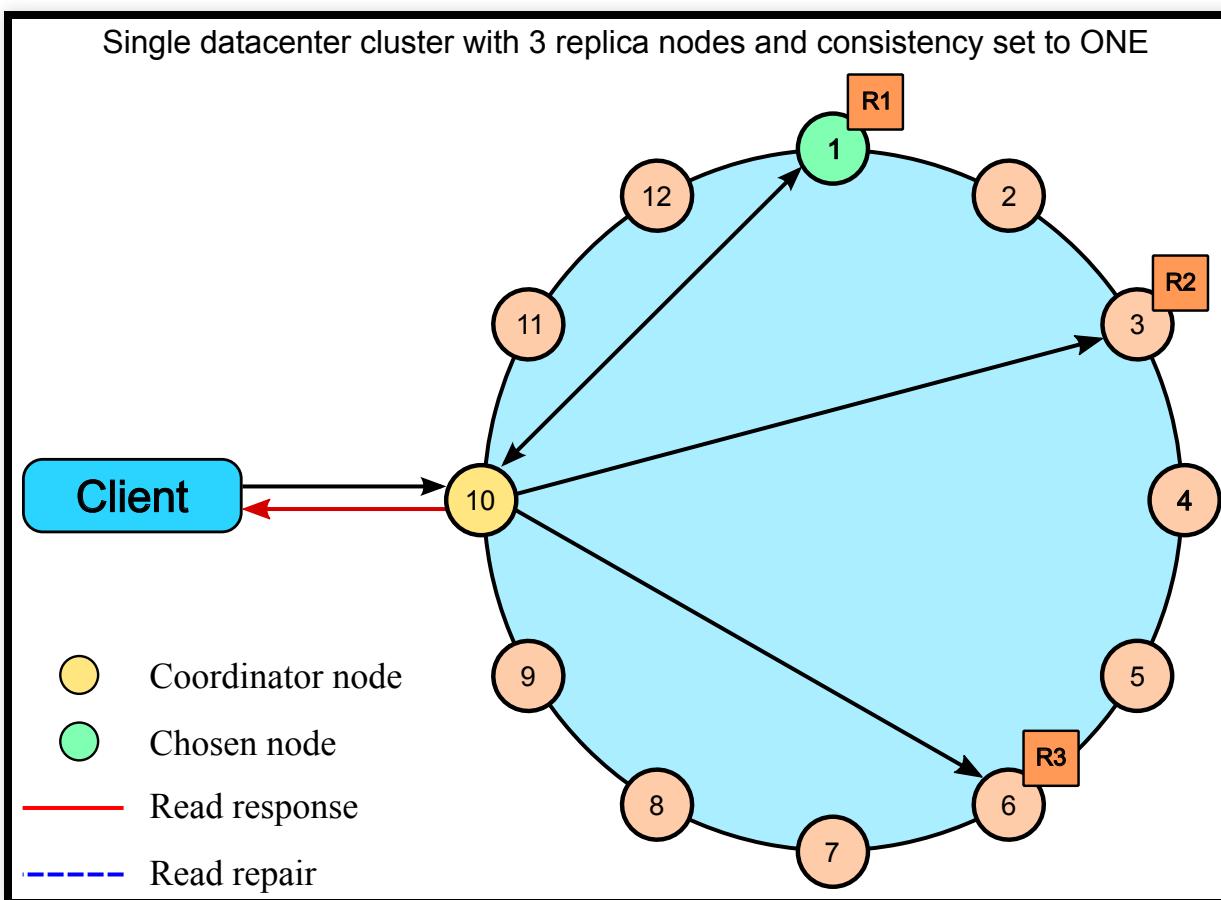
- ONE/QUORUM vs ANY (any node may ACK even if not a replica)
- if one/more replica(s) are down  $\Rightarrow$  hinted handoff



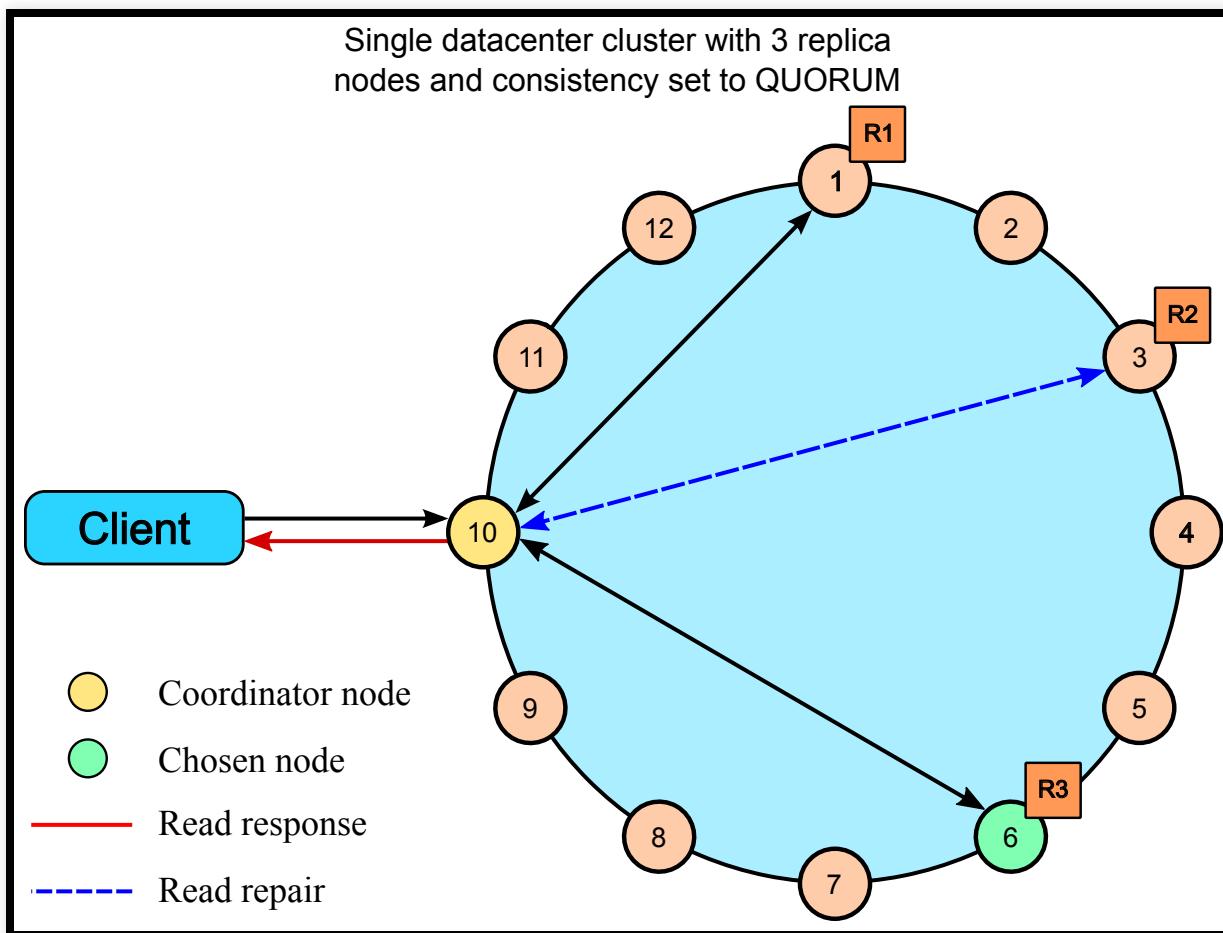
# Read repairs

- Goal: detect and fix inconsistencies during reads
- **CL = ONE**  $\Rightarrow$  no data is repaired because no comparison takes place (unless `read_repair_chance > 0`)
- **CL > ONE**  $\Rightarrow$  repair participating replica nodes in the foreground before the data is returned to the client.
  - C\* sends a digest request to each replica not directly involved in the read
  - Cassandra compares all replicas and writes the most recent version to any replica node that does not have it.

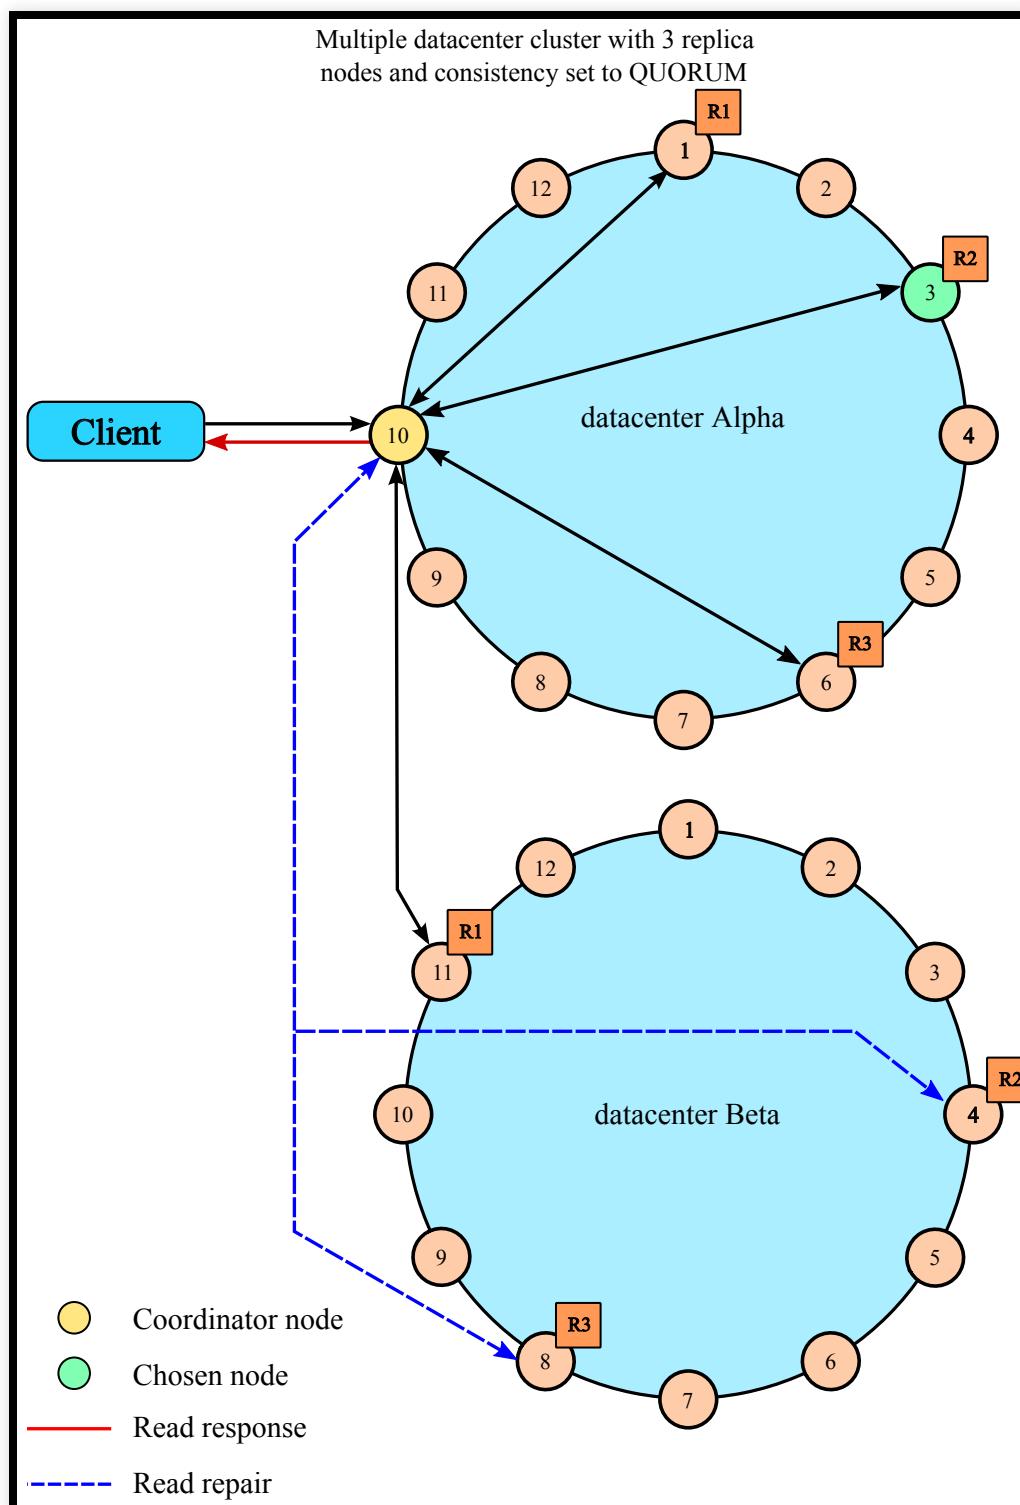
# Read repairs ONE



# Read repairs QUORUM



# Read repairs QUORUM DC



# Anti-entropy repair

- for each token range, read and synchronize the rows
- to insure the consistency this tool must be run regularly !
- not automatised (only on Datastax...)

# Plan

- Motivation
- Apache Cassandra
  - Partitioning and replication
  - Consistency
- **Practice: Tune consistency in Apache Cassandra**

# Practice: Tune consistency in Apache Cassandra

1. create local test clusters
2. explore configuration options and consistency properties

# Cassandra cluster manager

- create multi-node cassandra clusters on the local machine
- great for quickly setting up clusters for development and testing

```
$ccm create test -v 2.0.5 -n 3 -s (1)  
$ccm node1 stop (2)  
$ccm node1 cqlsh (3)
```

# Nodetool

- a command line interface for managing a cluster
  - explore, debug, performance test
  - maintenance operations, repairs

```
$ ccm node1 nodetool status mykeyspace (1)

Datacenter: datacenter1
=====
Status=Up/Down
| / State=Normal/Leaving/Joining/Moving
--  Address      Load      Tokens  Owns      Host ID
UN  127.0.0.1    47.66 KB     1       33.3%  aa1b7c1-6049-4a08-ad3e-3697a0e30e1
UN  127.0.0.2    47.67 KB     1       33.3%  1848c369-4306-4874-afdf-5c1e95b8732
UN  127.0.0.3    47.67 KB     1       33.3%  49578bf1-728f-438d-b1c1-d8dd644b6f7
```

# CQLSh

- standard CQL client

```
[bigdata@bigdata ~]$ ccm node2 cqlsh (1)
Connected to test at 127.0.0.2:9160.
[cqlsh 4.1.1 | Cassandra 2.0.5-SNAPSHOT | CQL spec 3.1.1 | Thrift protocol 19
Use HELP for help.
cqlsh> SELECT * FROM system.schema_keyspaces ; (2)
```

# Ressources:

Datastax documentation

<https://dzone.com/articles/introduction-apache-cassandras>

<https://highlyscalable.wordpress.com/2012/09/18/distributed-algorithms-in-nosql-databases/>